

Logging and Recovery

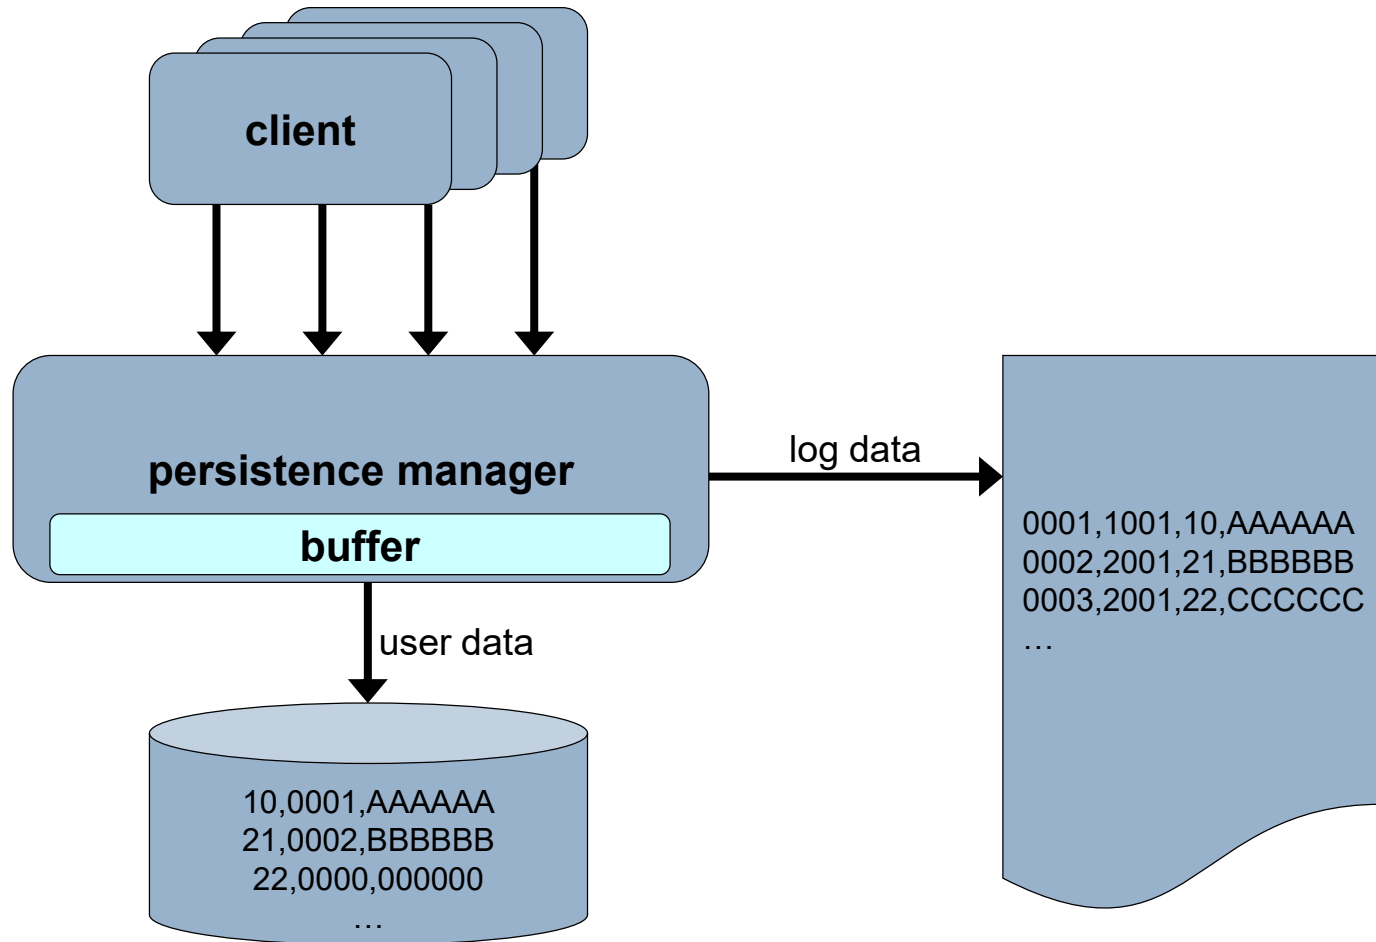
DIS Exercise Course



Task

- implementing a simplified „database system“
- several clients access one single persistence
- buffer management
 - Non-Atomic
 - No-Steal
 - No-Force
- logging during normal operation
 - physical state logging
- crash recovery
 - redo recovery only
 - re-issuing lost modifications

Architecture



Persistence Manager

- persisting pages
 - in files, not in DB2
 - advice: one file per page
- page structure: [PageID, LSN, Data]
 - PageID: Page identifier
 - LSN: Log Sequence Number
 - Data: user data
- buffer: delayed writes
 - persisting data of **committed** transactions, when there are more than five pages in the buffer (No-Force)
 - check for a „full“ buffer after every write access
 - refreshing pages in the buffer is possible
 - no dirty pages are persisted (No-Steal)
 - user data are overwritten directly (Non-Atomic)
- modifying operations are logged immediately
- requirement: support of crash recovery

Persistence Manager: Logging

- physical state logging
- granulate: Page
- new states (after-images) of modified objects are written to the log file
- record types
 - BOT and commit record
 - modification records (after-images)
- structure of log entries: [LSN, TAID, PageID, Redo]
 - LSN: Log Sequence Number (monotonically ascending)
 - TAID: transaction ID
 - PageID: Page ID
 - Redo: information required for redo

Clients

- 5 clients of the same kind
 - concurrent threads with ClientIDs 1..5
- access to persistence manager
 - persistence manager as Singleton: a single instance
 - clients write pages like this:
`beginTransaction() write() write()... commit()`
 - no concurrent access on the same object → no locks required
(Client 1: pages 10..19, Client 2: pages 20..29 etc.)
 - TAID is assigned by the persistence manager

Crash Recovery

- a record in the DB is either
 - up-to-date
 - or stale (noforce) → redo
- analysis phase
 - read the entire log sequentially
 - determine winner TAs
- redo phase
 - read the entire log sequentially
 - selective redo (redo winners)
- recovery progress is logged by updating the LSNs of the completed redo steps in the corresponding pages

Crash Recovery

- redo is only required, if
page LSN < LSN of redo log record
- page LSN is updated on redo (grows monotonically)

```
if LSN(LS) > LSN(Page) then  
    redo (modification from LS);  
    LSN(Page) := LSN(LS);  
end;
```